

Submodular Maximization

advances in distributed/streaming computing

Jiecao Chen

Indiana University Bloomington

jiēcchen@indiana

March 17, 2016

Overview

- 1 Introduction to Submodularity
 - Definitions
 - Properties
 - Constraints
 - Algorithms
- 2 Applications
 - Overview
 - List of Possible Applications
 - Examples of Applications
- 3 Streaming Submodular Maximization
 - Streaming Model
 - Algorithms
- 4 Distributed Submodular Maximization

Definitions of Submodularity

Definition (submodular concave)

A function $f : 2^V \rightarrow \mathbb{R}$ is **submodular** if for any $A, B \subseteq V$, we have that:

$$f(A) + f(B) \geq f(A \cup B) + f(A \cap B). \quad (1)$$

An alternate equivalent definition is more interpretable in many situations.

Definition (diminishing returns)

A function $f : 2^V \rightarrow \mathbb{R}$ is **submodular** if for any $A \subseteq B \subset V$, and $v \in V \setminus B$, we have that:

$$f(A + v) - f(A) \geq f(B + v) - f(B). \quad (2)$$

Modular Functions

Definition (Modularity)

A function $f : 2^V \rightarrow \mathbb{R}$ is **modular** if for any $A \subseteq B \subset V$, and $v \in V \setminus B$, we have that:

$$f(A + v) - f(A) = f(B + v) - f(B). \quad (3)$$

Notably, a modular function f can always be written as

$$f(S) = f(\emptyset) + \sum_{v \in S} (f(\{v\}) - f(\emptyset))$$

for any $S \subseteq V$. If we further assume $f(\emptyset) = 0$ (in this case, we call f **normalized** or **proper**), we have a simplified expression,

$$f(S) = \sum_{v \in S} f(\{v\}).$$

Monotonicity

Definition (Monotonicity)

A set function $f : 2^V \rightarrow \mathbb{R}$ is said to be non-decreasing if for any $A \subseteq B \subseteq V$, $f(A) \leq f(B)$. Non-increasing set functions are defined in the similar way.

When we say a submodular function is monotone, we mean it is non-decreasing.

Properties

Submodularity is closed under addition.

Property

Let $f_1, f_2 : 2^V \rightarrow \mathbb{R}$ be two submodular functions. Then

$$f : 2^V \rightarrow \mathbb{R} \text{ with } f(A) = \alpha f_1(A) + \beta f_2(A)$$

is submodular for any fixed $\alpha, \beta \in \mathbb{R}^+$.

Submodularity is preserved under restriction.

Property

Let $f : 2^V \rightarrow \mathbb{R}$ be a submodular function. Let $S \subseteq V$ be a fixed set. Then

$$f' : 2^V \rightarrow \mathbb{R} \text{ with } f'(A) = f(A \cap S)$$

is submodular.

Properties cont.

The following property can be useful when we show that the negative of the objective function of k-median problem is submodular.

Property

Consider V as a set of indices. Let $\mathbf{c} \in \mathbb{R}^V$ be a fixed vector, c_i its i th coordinate. Then

$$f : 2^V \rightarrow \mathbb{R} \text{ with } f(A) = \max_{j \in A} c_j$$

is submodular.

Constraints

Submodular Maximization Problem

A submodular maximization problem usually has the following form:

$$\arg \max_{I \in \mathcal{I}} f(I), \quad (4)$$

where f is a submodular function and $\mathcal{I} \subseteq 2^V$ is the collection of all feasible solutions. We call \mathcal{I} the **constraint** of the optimization problem.

Constraints

\mathcal{I} is important!

The structure of \mathcal{I} plays a crucial role in submodular optimization:

- Different constraints have different hardness results.
- Normally the difficulty increases when the constraint becomes more general.

Constraints

\mathcal{I} is important!

The structure of \mathcal{I} plays a crucial role in submodular optimization:

- Different constraints have different hardness results.
- Normally the difficulty increases when the constraint becomes more general.

Popular constraints

Some popular constraints:

- Cardinality constraint
- Knapsack constraint
- Matroid constraint
- Matching
- p -System
- ...

Constraints cont.

First we define hereditary set systems.

Definition (Hereditary)

A constraint $\mathcal{I} \subseteq 2^V$ is said to be **hereditary** if

$$I \in \mathcal{I} \implies J \in \mathcal{I} \text{ for any } J \subseteq I.$$

A hereditary constraint is sometimes called an **independent system** and each $I \in \mathcal{I}$ is called an **independent set**.

All constraints we will discuss are hereditary.

Constraints cont.

Cardinality

Cardinality constraint: $\mathcal{I} = \{A \subseteq V \mid |A| \leq k\}$

Knapsack

Knapsack Constraint: each $i \in V$ is assigned a weight $w_i \geq 0$,
 $\mathcal{I} = \{S \subseteq V \mid \sum_{i \in S} w_i \leq W\}.$

Constraints cont.

Cardinality

Cardinality constraint: $\mathcal{I} = \{A \subseteq V \mid |A| \leq k\}$

Knapsack

Knapsack Constraint: each $i \in V$ is assigned a weight $w_i \geq 0$,
 $\mathcal{I} = \{S \subseteq V \mid \sum_{i \in S} w_i \leq W\}$.

Matching

Matching: given a graph $G = (V, E)$, a *Matching* is a set $S \subseteq E$ such that no edges in S share common vertex.

Matroid

Matroid is the generalization of the independence concept in linear algebra; omit details here ...

p -System

p -system is very general, it includes many other constraints as special cases.

Definition of p -System

Let (V, \mathcal{I}) be a set system and \mathcal{I} hereditary. Let $\mathcal{B}(A)$ be the collection of all bases of A .

$$\mathcal{I} = \{A \subseteq V \mid \frac{\max_{S \in \mathcal{B}(A)} |S|}{\min_{S \in \mathcal{B}(A)} |S|} \leq p\}.$$

p -System

p -system is very general, it includes many other constraints as special cases.

Definition of p -System

Let (V, \mathcal{I}) be a set system and \mathcal{I} hereditary. Let $\mathcal{B}(A)$ be the collection of all bases of A .

$$\mathcal{I} = \{A \subseteq V \mid \frac{\max_{S \in \mathcal{B}(A)} |S|}{\min_{S \in \mathcal{B}(A)} |S|} \leq p\}.$$

Note: a base of A is the maximal independent set included in A .

p -System

p -system is very general, it includes many other constraints as special cases.

Definition of p -System

Let (V, \mathcal{I}) be a set system and \mathcal{I} hereditary. Let $\mathcal{B}(A)$ be the collection of all bases of A .


$$\mathcal{I} = \{A \subseteq V \mid \frac{\max_{S \in \mathcal{B}(A)} |S|}{\min_{S \in \mathcal{B}(A)} |S|} \leq p\}.$$

Note: a base of A is the maximal independent set included in A .

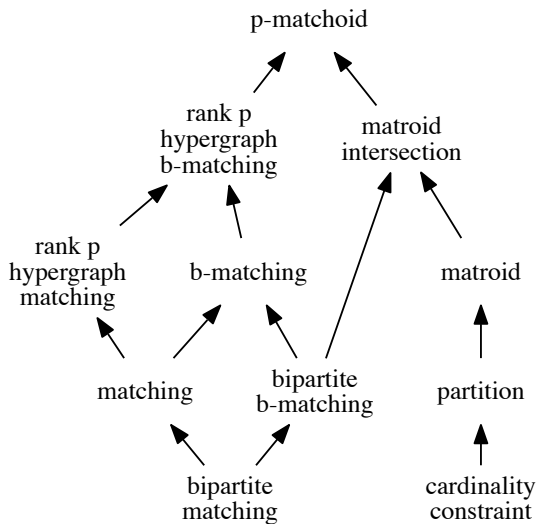
examples of p -system

- matroid is 1-system
- matching is 2-system
- intersection of p matroid is p -system
- ...

Hierarchy of constraints



Hierarchy of constraints (extended)



Notations

Some notations

- $\Delta_f(e|S) = f(S + e) - f(S)$ (or simply $\Delta(e|S)$ when f is clear from context)
- **α -approximation**: the returned solution S always satisfies $f(S) \geq \alpha \cdot \arg \max_{I \in \mathcal{I}} f(I)$
- When the algorithm is randomized, we normally say it guarantees **α -approximation in expectation** if

$$\mathbf{E}[f(S)] \geq \alpha \cdot \arg \max_{I \in \mathcal{I}} f(I).$$

The standard greedy algorithm

Algorithm 1: GREEDY algorithm for submodular maximization

Input: V the ground set, f the submodular function, \mathcal{I} the constraint

Output: a set $S \subseteq V$

```

1  $S \leftarrow \emptyset$ 
2 while  $\exists e \in V \setminus S$  s.t.  $S \cup \{e\} \in \mathcal{I}$  do
3    $e \leftarrow \arg \max_{e \in V \setminus S, S \cup \{e\} \in \mathcal{I}} \Delta_f(e|S)$ 
4    $S \leftarrow S \cup \{e\}$ 
5 return  $S$ 
```

Theorems of Algorithm 1

Theorem ([10], for cardinality constraint)

For a non-negative *monotone submodular* function $f : 2^V \rightarrow \mathbb{R}$, let \mathcal{I} be the *cardinality constraint*, Algorithm 1 returns a $(1 - 1/e)$ -approximation to $\arg \max_{I \in \mathcal{I}} f(S)$.

Theorems of Algorithm 1

Theorem ([10], for cardinality constraint)

For a non-negative *monotone submodular* function $f : 2^V \rightarrow \mathbb{R}$, let \mathcal{I} be the *cardinality constraint*, Algorithm 1 returns a $(1 - 1/e)$ -approximation to $\arg \max_{I \in \mathcal{I}} f(S)$.

Theorem ([10, 2], for p -system)

For a non-negative *monotone submodular* function f , given a p -system (V, \mathcal{I}) , Algorithm 1 returns a $\frac{1}{p+1}$ -approximation.

Theorems of Algorithm 1

Theorem ([10], for cardinality constraint)

For a non-negative *monotone submodular* function $f : 2^V \rightarrow \mathbb{R}$, let \mathcal{I} be the *cardinality constraint*, Algorithm 1 returns a $(1 - 1/e)$ -approximation to $\arg \max_{I \in \mathcal{I}} f(S)$.

Theorem ([10, 2], for p -system)

For a non-negative *monotone submodular* function f , given a p -system (V, \mathcal{I}) , Algorithm 1 returns a $\frac{1}{p+1}$ -approximation.

Theorem ([5], modular maximization s.t. p -system)

For a non-negative *monotone modular* function f , given a p -system (V, \mathcal{I}) , Algorithm 1 returns a $\frac{1}{p}$ -approximation.

Speedup - GREEDYLAZY

GreedyLazy

- Minoux [8] proposed LAZY-GREEDY as a fast implementation for Algorithm 1.

Speedup - GREEDYLAZY

GreedyLazy

- Minoux [8] proposed LAZY-GREEDY as a fast implementation for Algorithm 1.
- GREEDYLAZY keeps an upper bound $\rho(e)$ on the marginal gain sorted in a heap.

Speedup - GREEDYLAZY

GreedyLazy

- Minoux [8] proposed LAZY-GREEDY as a fast implementation for Algorithm 1.
- GREEDYLAZY keeps an upper bound $\rho(e)$ on the marginal gain sorted in a heap.
- In each step, only update the top element in the heap and push it back, if this element remains in the top, include it into solution.

Speedup - GREEDYLAZY

GreedyLazy

- Minoux [8] proposed LAZY-GREEDY as a fast implementation for Algorithm 1.
- GREEDYLAZY keeps an upper bound $\rho(e)$ on the marginal gain sorted in a heap.
- In each step, only update the top element in the heap and push it back, if this element remains in the top, include it into solution.
- Again gives $(1 - e^{-1})$ -approximation.

Speedup - STOCGREEDY[9]

StocGreedy

- In each round, instead of considering all $V \setminus S$ to get

$$e \leftarrow \arg \max_{e \in V \setminus S, S \cup \{e\} \in \mathcal{I}} \Delta_f(e|S),$$

Speedup - STOCGREEDY[9]

StocGreedy

- In each round, instead of considering all $V \setminus S$ to get

$$e \leftarrow \arg \max_{e \in V \setminus S, S \cup \{e\} \in \mathcal{I}} \Delta_f(e|S),$$

- consider only $\frac{|V|}{k} \log \frac{1}{\epsilon}$ random samples from $V \setminus S$.

Speedup - STOCGREEDY[9]


StocGreedy

- In each round, instead of considering all $V \setminus S$ to get


$$e \leftarrow \arg \max_{e \in V \setminus S, S \cup \{e\} \in \mathcal{I}} \Delta_f(e|S),$$

- consider only $\frac{|V|}{k} \log \frac{1}{\epsilon}$ random samples from $V \setminus S$.
- $(1 - e^{-1} - \epsilon)$ -approximation in expectation.

Comparison



(a) Efficiency



(b) Quality

Figure : Experiment on SYNTHETIC dataset

Summary of state of the arts

constraint	monotone	non-negative
cardinality	$1 - 1/e$ [10]	$1/e + .004$ [1]
matroid	$1 - 1/e$ [2], R	$\frac{1-\epsilon}{e}$ [3], R
matching	$\frac{1}{2+\epsilon}$ [4]	$\frac{1}{4+\epsilon}$ [4]
intersection of p matroids	$\frac{1}{p+\epsilon}$ [7]	$\frac{p-1}{p^2+\epsilon}$ [7]
p -matchoid	$\frac{1}{p+1}$ [2, 10]	$\frac{(1-\epsilon)(2-o(1))}{e \cdot p}$ [4, 11], R

Table : Best known approximation bounds for submodular maximization in RAM model. Bounds for randomized algorithms that hold in expectation are marked (R).

Overview of Applications

- **Combinatorial Problems:** set cover, max k coverage, vertex cover, edge cover, graph cut problems etc.
- **Networks:** social networks, viral marketing, diffusion networks etc.
- **Graphical Models:** image segmentation, tree distributions, factors etc.
- **NLP:** document summarization, web search, information retrieval
- **Machine Learning:** active/semi-supervised learning etc.
- **Economics:** markets, economies of scale

Set Cover Problem

- Let E be a fixed set with finite size.
- $V = \{C_1, C_2, \dots, C_n\}$ where each $C_i \subseteq E$.
- We define a function $f : 2^V \rightarrow \mathbb{R}$ such that $f(S) = |\cup_{C \in S} C|$.
- Goal: pick $S \subseteq V$ with $|S| \leq k$ that maximizes $f(S)$
- $f(S)$ is monotone submodular and this is a submodular maximization problem s.t. cardinality constraint!

Kernel Machines

The data set $V = \{x_1, x_2, \dots, x_n\}$ is represented in a transformed space via a kernel matrix

$$K_V = \begin{pmatrix} \mathcal{K}(x_1, x_1) & \dots & \mathcal{K}(x_1, x_n) \\ \vdots & \ddots & \vdots \\ \mathcal{K}(x_n, x_1) & \dots & \mathcal{K}(x_n, x_n) \end{pmatrix},$$

where $\mathcal{K} : V \times V \rightarrow \mathbb{R}$ is the kernel function that is symmetric and positive definite.

Kernel Machines cont.

- K_V is large for large $|V|$, need to select a subset from V .
- How to measure the quality of selected subset?
- A popular way is to use *Informative Vector Machine* (IVM) introduced by Laurence et al. [6]:

$$f(S) = \frac{1}{2} \log \det (\mathbf{I} + \sigma^{-2} K_S)$$

- $f(S)$ is submodular!
- Goal:

$$\arg \max_{S \subseteq V: |S| \leq k} f(S).$$

The model

The ground set V is an ordered sequence of items e_1, e_2, \dots, e_n .
We process the items one by one and the maximum space being used should be sublinear (i.e. $o(n)$).

CHEN : *insert a graph here*

SIEVESTREAM assume OPT is known

Algorithm 2: SIEVESTREAMOPT for submodular maximization

Input: V as data stream, f a monotone submodular function, k the size constraint, OPT the optimal value of $f(S)$ under the constraint

Output: a set $S \subseteq V$

```

1  $S \leftarrow \emptyset$ 
2 for each  $e$  in the data stream do
3   if  $\Delta(e|S) \geq \frac{\text{OPT}/2 - f(S)}{k - |S|}$  and  $|S| < k$  then
4      $S \leftarrow S \cup \{e\}$ 
5 return  $S$ 

```

SIEVESTREAM assume OPT is unknown

Problems with SIEVESTREAM_{OPT}

OPT is unknown!

So what we do?

Solution

- $m = \max_{e \in V} f(\{e\})$, for simplicity, assume $f(\emptyset) = \emptyset$
- note that $m \leq \text{OPT} \leq k \cdot m$
- if we know m , we guess OPT as $m, (1 + \epsilon)m, (1 + \epsilon)^2 m, \dots \leq k \cdot m$, each guess runs an instance of SIEVESTREAM_{OPT}
- it runs only $O(\log_{(1+\epsilon)} k) = O(\frac{k}{\epsilon})$ instances

SIEVESTREAM assume OPT is unknown, cont.

Problem again

calculating $m = \max_{e \in V} f(\{e\})$ requires an extra pass!

Solution?

Solution

- update $m \leftarrow \max(f(e_i, m))$ on the fly!
- lazy-evaluation, create an instance of SIEVESTREAMOPT only when necessary
- it runs only $O(\log_{(1+\epsilon)}) = O(\frac{k}{\epsilon})$ instances, using only 1 pass
- guarantee $(1/2 - \epsilon)$ -approximation for monotone submodular maximization s.t. cardinality constraint

RANDOMSTREAM

Algorithm 3: RANDOMSTREAM for submodular maximization

Input: V as data stream, f a non-negative submodular function, k the cardinality constraint, ϵ a parameter

Output: a set $S \subseteq V$

```

1  $B \leftarrow \emptyset, S \leftarrow \emptyset$ 
2 for each  $e$  in the data stream do
3   if  $|S| < k$  and  $\Delta(e|S) > \alpha$  then
4      $B \leftarrow B + e$ 
5   if  $|B| > \frac{k}{\epsilon}$  then
6      $e \leftarrow$  uniformly random from  $B$ 
7      $B \leftarrow B - e, S \leftarrow S + e$ 
8     for all  $e' \in B$  s.t.  $\Delta(e'|S) \leq \alpha$  do
9        $B \leftarrow B - e'$ 
10  $S' \leftarrow$  offline algorithm on  $B$ 
11 return  $\arg \max_{A \in \{S, S'\}} f(A)$ 

```

Thank you!



N. Buchbinder, M. Feldman, J. S. Naor, and R. Schwartz.

Submodular maximization with cardinality constraints.

In *Proceedings of the Twenty-Fifth Annual ACM-SIAM Symposium on Discrete Algorithms*, pages 1433–1452. SIAM, 2014.



G. Calinescu, C. Chekuri, M. Pál, and J. Vondrák.

Maximizing a monotone submodular function subject to a matroid constraint.

SIAM Journal on Computing, 40(6):1740–1766, 2011.



M. Feldman, J. Naor, and R. Schwartz.

A unified continuous greedy algorithm for submodular maximization.

In *Foundations of Computer Science (FOCS), 2011 IEEE 52nd Annual Symposium on*, pages 570–579. IEEE, 2011.



M. Feldman, J. S. Naor, R. Schwartz, and J. Ward.

Improved approximations for k-exchange systems.

In *Algorithms–ESA 2011*, pages 784–798. Springer, 2011.



T. A. Jenkyns.

The efficacy of the greedy algorithm.

In *Proceedings of the 7th Southeastern Conference on Combinatorics, Graph Theory, and Computing, Utilitas Mathematica, Winnipeg*, pages 341–350, 1976.



N. Lawrence, M. Seeger, and R. Herbrich.

Fast sparse gaussian process methods: The informative vector machine.

In *Proceedings of the 16th Annual Conference on Neural Information Processing Systems*, number EPFL-CONF-161319, pages 609–616, 2003.



J. Lee, M. Sviridenko, and J. Vondrák.

Submodular maximization over multiple matroids via generalized exchange properties.

Mathematics of Operations Research, 35(4):795–806, 2010.



M. Minoux.

Accelerated greedy algorithms for maximizing submodular set functions.

In *Optimization Techniques*, pages 234–243. Springer, 1978.



B. Mirzasoleiman, A. Badanidiyuru, A. Karbasi, J. Vondrak, and A. Krause.

Lazier than lazy greedy.

In *Twenty-Ninth AAAI Conference on Artificial Intelligence*, 2015.



G. L. Nemhauser, L. A. Wolsey, and M. L. Fisher.

An analysis of approximations for maximizing submodular set functions.

Mathematical Programming, 14(1):265–294, 1978.



J. Vondrák, C. Chekuri, and R. Zenklusen.

Submodular function maximization via the multilinear relaxation and contention resolution schemes.

In *Proceedings of the forty-third annual ACM symposium on Theory of computing*, pages 783–792. ACM, 2011.